

# RAPIDSMITH

---

## **A Library for Low-level Manipulation of Partially Placed-and-Routed FPGA Designs**

### **Technical Report and Documentation**

Christopher Lavin, Marc Padilla, Jaren Lamprecht, Philip Lundrigan,  
Brent Nelson, Brad Hutchings, and Michael Wirthlin

NSF Center for High Performance Reconfigurable Computing (CHREC)  
Department of Electrical and Computer Engineering  
Brigham Young University  
Provo, UT 84602

# CONTENTS

---

Introduction.....	5
What is RapidSmith? .....	5
Who Should Use RapidSmith? .....	5
Why RapidSmith? .....	5
Which Xilinx Parts does RapidSmith Support? .....	5
How is This Different than VPR? .....	6
Why Java? .....	6
Legal and Dependencies .....	7
RapidSmith Legal Text .....	7
Included Dependency Projects .....	7
Getting Started .....	9
Installation .....	9
Getting RapidSmith .....	9
Requirements for Installation .....	9
Steps for Installation .....	9
Additional Notes for Mac OS X Installation .....	10
Overview .....	10
design Package .....	11
design.explorer Package .....	13
design.parser Package .....	13
device Package .....	13
device.browser Package .....	14
device.helper Package .....	15
examples Package .....	15
primitiveDefs Package .....	15
placer Package .....	15
router Package .....	16
util Package .....	16
Examples .....	16
Hello World .....	16
Hand Router .....	18
Part Tile Browser .....	18
Understanding XDL .....	19

What is XDL? .....	19
Basic Syntax of XDL Files .....	20
Design Statement .....	20
Module Statement .....	20
Instance Statement .....	21
Net Statement.....	21
Basic Syntax of XDLRC Files .....	22
Tiles.....	22
Primitive Sites.....	22
Wire.....	23
PIP.....	23
Primitive Definitions.....	24
RapidSmith Structure.....	25
A RapidSmith Design .....	25
Loading Designs .....	25
Saving Designs.....	26
A RapidSmith Device .....	26
Device .....	26
Wire Enumerator.....	26
Memory and Performance.....	27
Virtex 4 Device Performance and Memory Usage .....	27
Virtex 5 Device Performance and Memory Usage .....	27
Virtex 6 Device Performance and Memory Usage .....	28
Spartan 6 Device Performance and Memory Usage .....	28
Wire Enumerator Size and Performance.....	29
Routing in RapidSmith .....	30
Wire Resources in RapidSmith.....	30
Wire Representation.....	30
Basic Routing.....	32
Router Structure .....	32
Routing Static Sources (VCC/GND) .....	33
Routing Clocks.....	33
Appendix.....	34
Appendix A: Modifying LUT Content .....	34
LUT Equation Syntax .....	34
XDL LUT Equation Syntax .....	34

Appendix B: Hard Macros in XDL and RapidSmith..... 35

    Xilinx NMC files ..... 35

    RapidSmith Hard Macro Generator ..... 35

# INTRODUCTION

---

## ***What is RapidSmith?***

The BYU RapidSmith project is a Java-based API that aims to provide academics with an *easy-to-use* platform to try out experimental ideas and algorithms on modern Xilinx FPGAs. RapidSmith is based on the Xilinx Design Language (XDL) which provides a human-readable file format equivalent to the Xilinx proprietary Netlist Circuit Description (NCD). With RapidSmith, researchers are able to import XDL/NCD, manipulate, place, route and export designs among a variety of possible design transformations possible. The RapidSmith project makes an excellent test bed to try out new ideas and algorithms for FPGA CAD research as code can quickly be written to take advantage of the APIs available.

The RapidSmith project does not manipulate Xilinx bitstreams and does not contain any information of the organization of such files. It also does not contain or provide methods of including FPGA timing information. RapidSmith does not include any proprietary information about Xilinx FPGAs that is not publicly available.

## ***Who Should Use RapidSmith?***

RapidSmith is aimed at use by academics in all fields of FPGA CAD research. It is written in Java; therefore those using it will need to have a basic knowledge of programming and using Java. It also depends on some understanding of Xilinx FPGAs and XDL, however, this documentation hopes to bring people unfamiliar with these topics up to speed.

RapidSmith by no means is a Xilinx ISE replacement and **cannot** be used without a valid and current license to a Xilinx tools installation. RapidSmith should not be used for designs bound for commercial products and is offered mainly as a research tool.

## ***Why RapidSmith?***

The Xilinx ISE tools provide an `xdl` executable that allows conversion of NCD files to and from XDL which can then be parsed, manipulated and exported using RapidSmith. The `xdl` executable also creates special device files which are huge in size but contain useful detailed device data.

RapidSmith takes care of all of the parsing and detailed FPGA part information that can be cumbersome to use—alleviating the need to build such parsing tools by the researcher. RapidSmith creates special part files from these device files created by the ISE tools which can then be used by RapidSmith for design manipulation. This project provides researchers the ability to leverage all of the XDL work previously done and avoid duplicate work. This will enable researchers to have more time to focus on what matters most: their research of new ideas and algorithms.

## ***Which Xilinx Parts does RapidSmith Support?***

Virtex 4 and 5 families have been tested the most and are currently supported. Portions of RapidSmith work with Virtex 6 and Spartan 6 families, however, they are not as extensively supported as Virtex 4 and 5. Despite these current limitations, the code is organized and written to support almost all FPGA families supported by ISE 11.1 to 12.2, these being: Virtex 4, Virtex 5, Virtex 6, Spartan 3A, Spartan 3ADSP, Spartan 3E and Spartan

6. There is currently no plans to support the Spartan 3 family, however, the code could be taken independently and changed to make it to work with the older architecture.

## ***How is This Different than VPR?***

[VPR \(Versatile Place and Route\)](#) has been an FPGA research tool for several years and has led to hundreds of publications on new FPGA CAD research. It has been a significant contribution to the FPGA research community and has grown to be a complete FPGA CAD flow for research-based FPGAs.

The main difference between RapidSmith and VPR is that RapidSmith aims to provide the ability to target commercial Xilinx FPGAs. All features of these FPGAs which are accessible via XDL are available in RapidSmith. Our understanding is that VPR currently is limited to FPGA features which can be described using VPR's architectural description facilities.

## ***Why Java?***

We have found Java to be a fast prototyping platform for FPGA CAD tools. The Java libraries are rich with data structures useful for such applications and Java eliminates the need to clean up objects in memory. This eliminates the time needed to debug such things in other development platforms, leaving more time for the researcher to focus on the real research at hand.

Some may argue that Java is a poor platform for FPGA CAD tool design as it has a reputation of being a memory hog and slow. We believe that these claims are overstated and that both speed and memory can be controlled to the point where this is not an issue.

# LEGAL AND DEPENDENCIES

---

## *RapidSmith Legal Text*

BYU RapidSmith Tools

Copyright (c) 2010 Brigham Young University

BYU RapidSmith Tools is free software: you may redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 2 of the License, or (at your option) any later version.

BYU RapidSmith Tools is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

A copy of the GNU General Public License is included with the BYU RapidSmith Tools. It can be found at doc/gpl2.txt. You may also get a copy of the license at <<http://www.gnu.org/licenses/>>.

## *Included Dependency Projects*

RapidSmith includes the Caucho Technology Hessian implementation which is under the Apache License which a copy of this license is included in this directory in APACHE2-LICENSE.txt. This license is also available for download at:

<http://www.apache.org/licenses/LICENSE-2.0>

The source for the Caucho Technology Hessian implementation is available at:

<http://hessian.caucho.com>

RapidSmith also includes the Qt Jambi project jars for Windows, Linux and Mac OS X. Qt Jambi is distributed under the LGPL GPL3 license and copies of this license and exception are also available in the /doc directory in files LICENSE.GPL3.TXT and LICENSE.LGPL.TXT respectively. These licenses can also be downloaded at:

<http://www.gnu.org/licenses/licenses.html>

Source for the Qt Jambi project is available at:

<http://qt.nokia.com/downloads>

and more recent versions are available at:

<http://qt.gitorious.org/qt-jambi>

The user is responsible for providing copies of these licenses and making available the source code of these projects when redistributing these jars with RapidSmith.



# GETTING STARTED

---

## Installation

### Getting RapidSmith

You can download the latest snapshot of the RapidSmith SVN repository which can be downloaded as a ZIP file from the sourceforge page here:

<https://sourceforge.net/projects/rapidsmith>

You can also checkout the repository from SVN. We recommend using Eclipse, however, any IDE will work fine. To check out the RapidSmith project, the SVN URL is:

<https://rapidsmith.svn.sourceforge.net/svnroot/rapidsmith> rapidsmith

Where you will want only the folder 'trunk'. This repository contains all the files you need (including supporting JAR files). If you are using Eclipse as your IDE, it contains project files to get the project up and running with minimal effort.

### Requirements for Installation

- 200 MB free disk space (for all Virtex 4 and Virtex 5 family devices)
  - NOTE: ~30-50 GB of free hard disk space needed during installation (XDLRC files are very large in size, the largest Virtex 6 parts take over 23GBs).
- Windows XP/Vista/7 or Linux (Mac OS X will work, but Xilinx tools do not so we currently don't support it)
- [Xilinx ISE](#) 11.1 or higher
- [JDK 1.6](#) (earlier versions may work, but have not been tested).
- Supporting JARs
  - INCLUDED: [Caucho Hessian Implementation JAR v.4.0.6](#) (Used for compressing database device files)
  - INCLUDED: [Qt Jambi](#) (Qt for Java) for the Part Tile Browser example. Just adding the [jars](#) to the CLASSPATH variable is adequate.
  - OPTIONAL: [JavaCC](#) if the user wants to change the XDL design parser. There is also a [good plugin for Eclipse for JavaCC](#) which makes it easier to modify and compile .jj files.

### Steps for Installation

1. Make sure the Xilinx tools and JDK are on your PATH.
2. Add the hessian-4.0.6.jar file to your CLASSPATH environment variable and Qt Jambi jars as well.
3. Add the RapidSmith Java project to your CLASSPATH environment variable.
4. Create an environment variable called RAPIDSMITH\_PATH and set its value to the path to where you have the Java project installed.
5. Compile all of the Java classes (this can be done automatically if the project is imported into an IDE such as Eclipse).
6. Generate the supporting device and enumeration files needed to run the various parts of RapidSmith. Please note that if you are generating both families of Virtex 4 and Virtex 5 parts, it will take **several**

**hours** and is best left to run **overnight** because of the time requirement. This only needs to be done once, however. To generate the part files, follow these steps:

- Choose which parts you plan to use, or you can choose to do all parts in the Virtex 4 and Virtex 5 families (in the future, more parts will be compatible).
- Run the installer for RapidSmith by running the main method in the class `edu.byu.ece.rapidSmith.util.Installer` by running the following at the command line:

```
Java -Xmx1600M edu.byu.ece.rapidSmith.util.Installer virtex4 virtex5
```

- The previous command will take several hours. Some of the larger parts will also require a lot of heap memory to generate the part file.
- You can test if the file generation worked by looking in the appropriate folders (`devices/virtex4` and `devices/virtex5`). You can also run the `BrowseDevice` class as a test to see if you are able to browse any of the parts that have just been created. You can run this with the following command:

```
Java edu.byu.ece.rapidSmith.util.BrowseDevice xc5vlx20tff323
```

## Additional Notes for Mac OS X Installation

- The Xilinx tools do not run under Mac OS X and therefore, the installer would have to generate the files on Linux or Windows. However, once the files are created, they could be moved to a Mac OS X installation.
- To add/modify global environment variables in Mac OS X, one preferred way would be to edit the `environment.plist` file in the `~/MacOSX` directory. Here is an example of a proper setup for RapidSmith:

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN"
"http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
  <key>RAPIDSMITH_PATH</key>
  <string>/Users/[user name]/Documents/workspace/rapidSmith</string>
  <key>CLASSPATH</key>
  <string>${CLASSPATH}:/Users/[user name]/Documents/workspace/rapidSmith:/Users/[user
name]/Documents/workspace/rapidSmith:/Users/[user
name]/Documents/workspace/rapidSmith/jars/hessian-4.0.6.jar:/Users/[user
name]/Documents/workspace/rapidSmith/jars/qtjambi-4.6.3.jar:/Users/[user
name]/Documents/workspace/rapidSmith/jars/qtjambi-macosx-gcc-4.6.3.jar</string>
</dict>
</plist>
```

- Another difference is that when running programs that use Qt in Java under Mac OS X, the user will need to supply an extra JVM switch, “`-XstartOnFirstThread`”

## Overview

RapidSmith is organized into several packages (all packages are prefixed with “`edu.byu.ece.rapidSmith`”):

Package Name	Description
<b>design</b>	Represents all of the constructs in XDL files (Instances, Nets, PIPs, Modules, and Designs).
<b>design.explorer</b>	This is a GUI interactive explorer that allows the user to navigate through the various constructs in the design (Nets, Instances, Modules and Module Instances). It also has a

	tile map which allows the user to view the locations of various objects on the FPGA fabric.
<b>design.parser</b>	A JavaCC-based parser for XDL files which populate an instance of the Design class in the design package.
<b>device</b>	This package encompasses all the details of an FPGA device (part name, tiles, primitive sites, routing resources). All information about Xilinx parts is populated in device from the XDLRC files generated by the xdl executable.
<b>device.browser</b>	This program is an extension of the part tile browser found in the examples package. It allows the user to browse all of the installed parts and also navigate primitive sites as well as routing resources.
<b>device.helper</b>	Some classes to help in the creation of the device files.
<b>examples</b>	Some user examples of how to use RapidSmith.
<b>primitiveDefs</b>	This is also populated from the XDLRC file, it is specific to a Xilinx family of parts (such as Virtex 4 or Virtex 5). It defines all primitives which are part of a Xilinx family of parts (SLICEL, SLICEM, RAMB16, ...).
<b>placer</b>	This contains classes to place designs.
<b>router</b>	This contains classes to route designs.
<b>util</b>	This contains miscellaneous support classes and utilities.

## design Package

The design package has all the essential classes necessary to represent all kinds of XDL designs with classes to represent each type of XDL construct. Below in Figures 1, 2, and 3 are some basic illustrations of how the most common XDL constructs map into RapidSmith design classes:



Figure 1 - Design and Attribute Classes

```
# =====
# The syntax for instances is:
#   instance <name> <sitedef>, placed <tile> <site>, cfg <string> ;
# or
#   instance <name> <sitedef>, unplaced, cfg <string> ;
#
# For typing convenience you can abbreviate instance to inst.
#
# For IOs there are two special keywords: bonded and unbonded
# that can be used to designate whether the PAD of an unplaced IO is
# bonded out. If neither keyword is specified, bonded is assumed.
#
# The bonding of placed IOs is determined by the site they are placed in.
#
# If you specify bonded or unbonded for an instance that is not an
# IOB it is ignored.
#
# Shown below are three examples for IOs.
#   instance IO1 IOB, unplaced ;           # This will be bonded
#   instance IO1 IOB, unplaced bonded ;    # This will be bonded
#   instance IO1 IOB, unplaced unbonded ;  # This will be unbonded
# =====
inst "counter<18>" "SLICEL", placed CLB X17Y51 SLICE X27Y103 ,
cfg " BXINV::#OFF BYINV::#OFF CEINV::#OFF CLKINV::CLK COUTUSED::0
CY0F::0 CY0G::0 CYINIT::CIN DXMUX::XMUX DYMUX::YMUX
F:counter<18> rt:#LUT:D=A1 _BEL_PROP::F:PK_PACKTHRU: F5USED::#OFF
FFX:counter_18:#FF FFX_INIT_ATTR::INIT0 FFX_SR_ATTR::SRLOW
FFY:counter_19:#FF FFY_INIT_ATTR::INIT0 FFY_SR_ATTR::SRLOW
FXMUX::FXOR FXUSED::#OFF G:counter<19>_rt:#LUT:D=A1
_BEL_PROP::G:PK_PACKTHRU: GYMUX::GXOR REVUSED::#OFF SRINV::#OFF
SYNC_ATTR::ASYNX XBUSED::#OFF XMUXUSED::#OFF XUSED::#OFF
YBUSED::#OFF YMUXUSED::#OFF YUSED::#OFF
CYMUXF:Mcount_counter_cy<18>: CYMUXG:Mcount_counter_cy<19>:
GNDG:ProtoCompl.GNDG.8: GNDG:ProtoCompl.GNDG.8:
XORF:Mcount_counter_xor<18>:
XORG:Mcount_counter_xor<19>:
_INST_PROP::XDL_SHAPE_MEMBER:Shape_0:0,9 "
```

## DESIGN CLASS

HashMap<String,Instance> instances;

## INSTANCE CLASS

- String name;
- PrimitiveType type;
- PrimitiveSite site;
- ArrayList<Attribute> attributes;

## ATTRIBUTE CLASS

String physicalName;  
String logicalName;  
String value;

Figure 2 - Instance Class

```
# Wire indicates that this a normal wire.
# Power indicates that this net is tied to a DC power source.
# You can use "power", "vcc" or "vdd" to specify a power net.
#
# Ground indicates that this net is tied to ground.
# You can use "ground", or "gnd" to specify a ground net.
#
# The <dir> token will be one of the following:
#
#   Symbol Description
#   =====
#   == Bidirectional, unbuffered.
#   => Bidirectional, buffered in one direction.
#   == Bidirectional, buffered in both directions.
#   -> Directional, buffered.
#
# No pips exist for unrouted nets.
# =====
net "GLOBAL LOGIC0_11" gnd,
  outpin "XDL_DUMMY_INT_X33Y26_TIEOFF_X33Y26" HARD0,
  inpin "FFT_F/U_CNTRL/Madd_ADDRD_addsub0000_cy<1>" BX,
  pip CLB X33Y26 BYP_INT B5_INT -> BX PINWIRE1,
  pip INT_X33Y26 BOUNCE3 -> BYP_INT B5,
  pip INT_X33Y26 GND_WIRE -> BOUNCE3,
;
```

## DESIGN CLASS

HashMap<String,Instance> instances;

## NET CLASS

- String name;
- NetType type;
- ArrayList<Pin> pins;

## PIN CLASS

- boolean isOutputPin;
- Instance instance;
- String pinName;
- ArrayList<PIP> pips;

## PIP CLASS

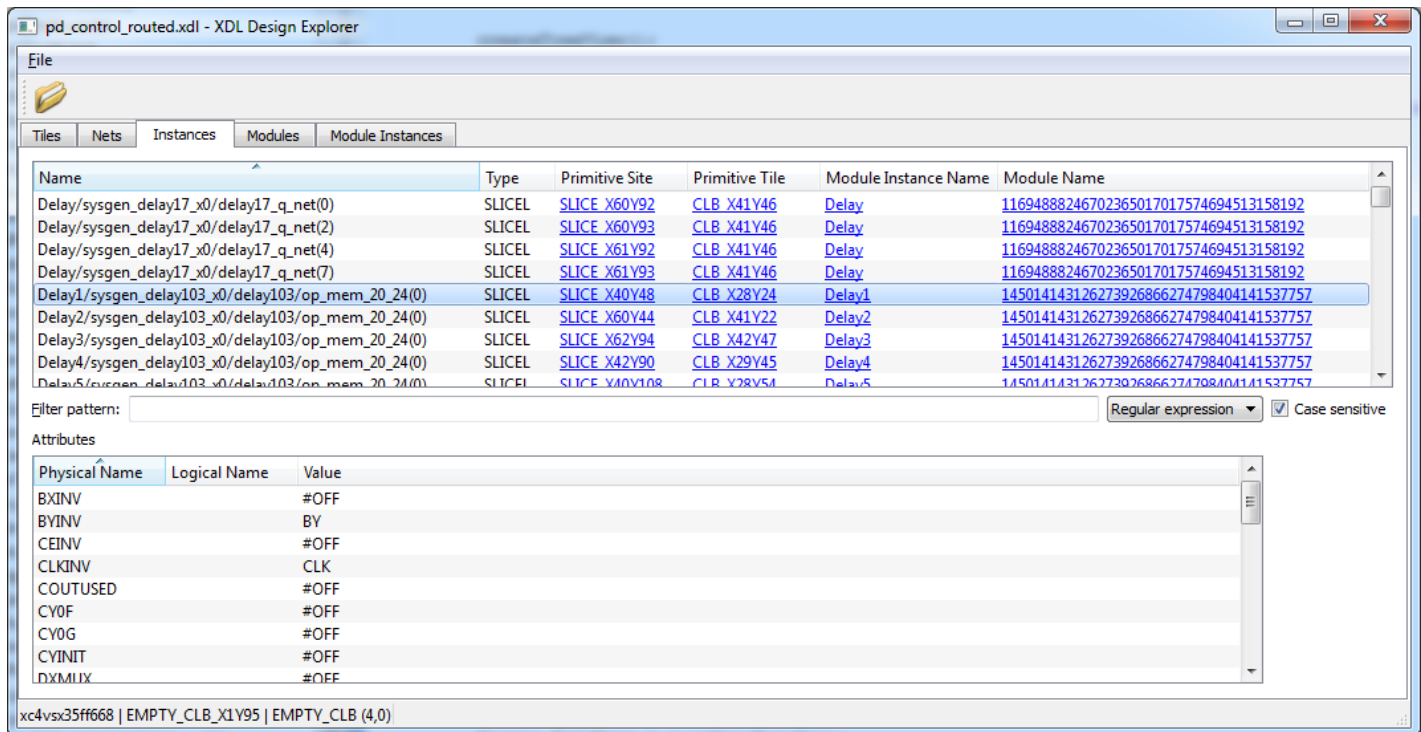
- Tile tile;
- int startWire;
- int endWire;

Figure 3 - Net, Pin and PIP Classes

There are other classes such as Module and ModuleInstance classes that abstract the macro-like property of XDL which will be explained later. There are also enumeration classes such as InstanceType which are an exhaustive list of all primitive types found in XDL and NetType which determines if a net is a WIRE, GND, or VCC.

## design.explorer Package

The design explorer loads XDL design files and allows the user a GUI interface to the design constructs with hyperlinks to the various other sites, tiles, modules and instances. Below is a screenshot from the design explorer.



## design.parser Package

The XDL design parser is written in JavaCC which compiles a .jj file into multiple .java files which implement a full-fledged parser. It will populate an instance of the Design class with the XDL design found in the .xdl file.

## device Package

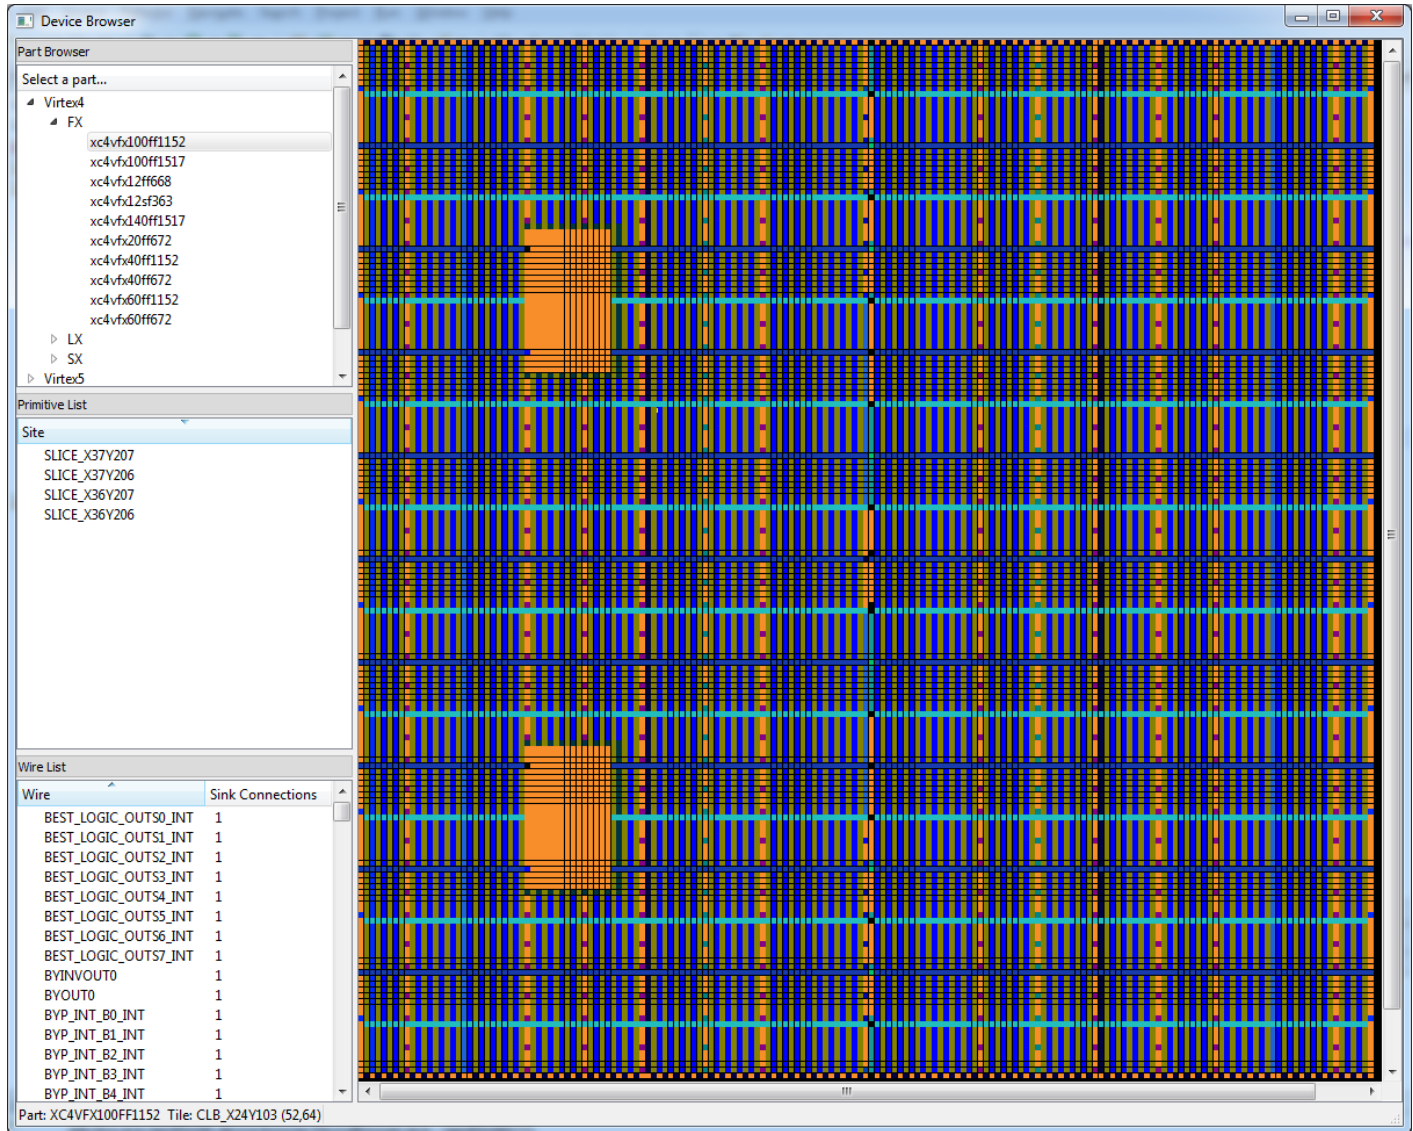
This package works closely with the design package in that the specific Device class is loaded when a design is loaded. The Xilinx XDLRC part descriptions partition the FPGA into a 2D grid of tiles. Each tile contains some mixture of primitive sites, wires and PIPs (Programmable Interconnect Points). Primitive sites are resource locations where XDL “inst” or instances of primitives are allowed to reside. Wires and PIPs provide wiring and routing resources information to connect the primitive instances together to form a complete design. With this information provided by Xilinx and leveraged by RapidSmith a number of different placement and routing algorithms can be constructed by leveraging the APIs in this package.

The device package also contains a class called WireEnumerator. All of the wires in a family are enumerated to an integer so they do not need to be stored as Strings. The WireEnumerator class helps translates wires from

integers to Strings and vice versa. It also keeps track of important information about wires such as the type of wire (DOUBLE, HEX, PENT, ...) wire direction (NORTH, SOUTH, EAST, ...) among other attributes.


## device.browser Package

The device browser application will allow the user to see a color coded tile array that allows them to browse any installed device. The primitive and wire lists are populated by double clicking a tile. The user can zoom in and out using the mouse wheel and can also pan by holding down the right mouse button while moving the mouse. See below for a screenshot of the application.



The device browser also allows the user to follow the various connections found in the FPGA. By double clicking a wire in the wire list, the application will draw the connection on the tile array (as shown in the screenshot below). By hovering the mouse pointer over the connection, the wire becomes red and a tooltip will appear describing the connection made by declaring the source tile and wire followed by an arrow (->) and the destination tile and wire. By clicking on the wire, the application will redraw all the connections that can be made from the currently selected wire. By repeating this action, the user can follow connections and discover how the FPGA interconnect is laid out.





### **device.helper Package**

This package contains special classes to help pack the device files smaller. They are generally used only during installation.

### **examples Package**

This package contains some examples of how to get started with RapidSmith and some different ways of using the various APIs available.

### **primitiveDefs Package**

In the XDLRC descriptions produced by the Xilinx 'xdl' executable, each copy has a section at the end called primitive\_defs which has a list of primitive definitions for all types of primitives found in the part. The primitiveDefs packages makes that information available in a convenient data structure to access the attributes and various parameters the primitives can be configured with.

### **placer Package**

This package still has yet to be completed but will have an example of a placer.

## router Package

This package still has yet to be completed but will have an example of a router that routes Virtex 4 and Virtex 5 designs.

## util Package

This has miscellaneous classes used for support of all other packages. It is suggested to have the user browse the JavaDoc API descriptions to get a better feel for what is contained in the util package.

# Examples

## Hello World

To get started programming with RapidSmith, here is an example of a very simple program.

```

/*
 * Copyright (c) 2010 Brigham Young University
 *
 * This file is part of the BYU RapidSmith Tools.
 *
 * BYU RapidSmith Tools is free software: you may redistribute it
 * and/or modify it under the terms of the GNU General Public License
 * as published by the Free Software Foundation, either version 2 of
 * the License, or (at your option) any later version.
 *
 * BYU RapidSmith Tools is distributed in the hope that it will be
 * useful, but WITHOUT ANY WARRANTY; without even the implied warranty
 * of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU
 * General Public License for more details.
 *
 * A copy of the GNU General Public License is included with the BYU
 * RapidSmith Tools. It can be found at doc/gpl2.txt. You may also
 * get a copy of the license at <http://www.gnu.org/licenses/>.
 */
package edu.byu.ece.rapidSmith.examples;

import java.util.HashMap;
import edu.byu.ece.rapidSmith.design.*;
import edu.byu.ece.rapidSmith.device.*;

/**
 * A simple class to illustrate how to use some of the basic methods in RapidSmith.
 * @author Chris Lavin
 */
public class HelloWorld{
    public static void main(String[] args){
        // Create a new Design from scratch rather than load an existing design
        Design design = new Design();

        // Set its name
        design.setName("helloWorld");

        // When we set the part name, it loads the corresponding Device and
        // WireEnumerator. Always include package and speed grade with the part name.
        design.setPartName("xc4vfx12ff668-10");

        // Create a new instance
        Instance myInstance = new Instance();
        myInstance.setName("Bob");
        myInstance.setType(PrimitiveType.SLICEL);
        // We need to add the instance to the design so it knows about it
        design.addInstance(myInstance);
        // Make the F LUT an Inverter Gate
    }
}

```



```

myInstance.addAttribute(new Attribute("F", "LUT_of_Bob", "#LUT:D=~A1"));

// Add the instance to the design
design.addInstance(myInstance);

// This is how we can get the reference to the instance from the design,
// by name
Instance bob = design.getInstance("Bob");

// Let's find a primitive site for our instance Bob
HashMap<String, PrimitiveSite> primitiveSites =
    design.getDevice().getPrimitiveSites();
for(PrimitiveSite site : primitiveSites.values()){
    // Some primitive sites can have more than one type reside at the site, such as
    // SLICEM sites which can also have SLICELs placed there. Checking if the site
    // is compatible makes sure you get the best possible chance of finding a place
    // for bob to live.
    if(site.isCompatiblePrimitiveType(bob.getType())){
        // Let's also make sure we don't place bob on a site that is already used
        if(!design.isPrimitiveSiteUsed(site)){
            bob.place(site);
            System.out.println("We placed bob on tile: " + bob.getTile() +
                               " and site: " + bob.getPrimitiveSiteName());
            break;
        }
    }
}

// Another way to find valid primitive sites if we want to use an exclusive site type
PrimitiveSite[] allSitesOfTypeSLICEL =
    design.getDevice().getAllSitesOfType(bob.getType());
for(PrimitiveSite site : allSitesOfTypeSLICEL){
    // Let's also make sure we don't place bob on a site that is already used
    if(!design.isPrimitiveSiteUsed(site)){
        bob.place(site);
        System.out.println("We placed bob on tile: " + bob.getTile() +
                           " and site: " + bob.getPrimitiveSiteName());
        break;
    }
}

// Let's create an IOB to drive our Inverter gate in Bob's LUT
Instance myIOB = new Instance();
myIOB.setName("input");
myIOB.setType(PrimitiveType.IOB);
design.addInstance(myIOB);
// These are typical attributes that need to be set to configure the IOB
// the way you like it
myIOB.addAttribute(new Attribute("INBUFUSED", "", "0"));
myIOB.addAttribute(new Attribute("IOATTRBOX", "", "LVCMOS25"));
// Another way to find a primitive site is by name, this is the pin name
// that you might find in a UCF file
myIOB.place(design.getDevice().getPrimitiveSite("C17"));

// Let's also create a new net to connect the two pins
Net fred = new Net();
// Be sure to add fred to the design
design.addNet(fred);
fred.setName("fred");
// All nets are normally of type WIRE, however, some are also GND and VCC
fred.setType(NetType.WIRE);
// Add the IOB pin as an output pin or the source of the net
fred.addPin(new Pin(true, "I", myIOB));
// Add Bob as the input pin or sink, which is the input to the inverter
fred.addPin(new Pin(false, "F1", bob));

// Now let's write out our new design

```

```

// We'll print the standard XDL comments out
String fileName = design.getName() + ".xdl";
design.saveXDLFile(fileName, true);

// We can load XDL files the same way.
Design inputFromFile = new Design();
inputFromFile.loadXDLFile(fileName);

// Hello World
System.out.println(inputFromFile.getName());
}
}

```

## Hand Router

This is a command-line-based router than allows a user to route one net at a time and write out the design changes afterwards. Although not particularly useful as a router, it illustrates how RapidSmith could be used to build a router.

## Part Tile Browser

This GUI will let you browse Virtex 4 and 5 parts at the tile level. On the left, the user may choose the desired part by navigating the tree menu and double-clicking on the desired part name. This will load the part in the viewer pane on the right (the first available part is loaded at startup). The status bar in the bottom left displays which part is currently loaded. Also displayed is the name of the current tile which the mouse is over, highlighted by a yellow outline in the viewer pane. The user may navigate inside the viewer pane by using the mouse. By right-clicking and dragging the cursor, the user may pan. By using the scroll-wheel on the mouse, the user may zoom. If a scroll-wheel is unavailable, the user may zoom by clicking inside the viewer pane and pressing the minus(-) key to zoom out or the equals(=) key to zoom in. See below for a screenshot.




Figure 4: Screenshot of Part Tile Browser

# UNDERSTANDING XDL

## What is XDL?

XDL (or Xilinx Design Language, see ISE 6.1 documentation in `help/data/xdl` folder) is a human-readable ASCII format compatible with the more widely used Xilinx format NCD (or Netlist Circuit Description). XDL has most if not all the same capabilities of the NCD format and Xilinx provides an executable called `xdl` which can convert NCD designs to XDL and vice versa (run “`xdl -h`” for details). XDL and NCD are both native Xilinx netlist formats for describing and representing Xilinx FPGA designs. XDL is the interface used by RapidSmith to insert and extract design information at different points in the Xilinx design flow.

XDL can represent designs that are:

- Mapped (unplaced and unrouted)
- Partially placed and unrouted
- Partially placed and partially routed
- Fully placed and unrouted
- Fully placed and partially routed
- Fully placed and fully routed
- Contain hard macros and instances of hard macros
- A hard macro definition (equivalent to Xilinx NMC files)

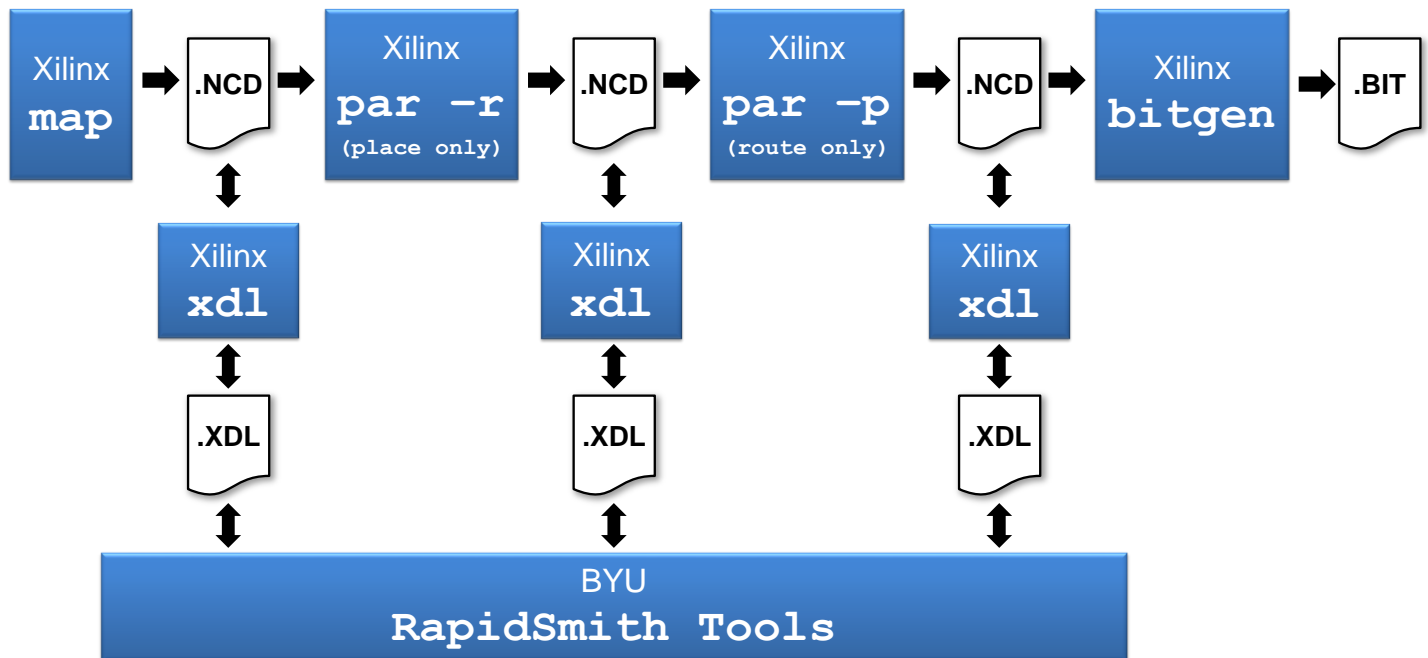


Figure 5: Block diagram of where XDL fits in CAD flow.

RapidSmith provides some Java methods that can perform the XDL/NCD conversion (by calling the `xdl` executable) from within Java in the `util.FileConverter` class. It also has method for calling a number of Xilinx programs from within the RapidSmith environment.

The Xilinx `xdl` executable also contains options for generating report files (with extension `.XDLRC`) which contain descriptive information about a particular Xilinx part. XDLRC report files have a different format to that of XDL (as they describe an FPGA rather than a design) and depending on the options given can create enormous files (several gigabytes) of text but are quite complete in describing the primitive sites, routing resources and tile layout of a Xilinx FPGA. RapidSmith makes use of these XDLRC files by generating them and parsing them into much smaller device files that can then be used with the rest of the RapidSmith API.

**DISCLAIMER:** The user must be aware that XDL is an externally **unsupported** format by Xilinx. All questions about XDL and any problems associated with XDL or this tool should **NOT** be addressed to Xilinx, but through the RapidSmith [website](#) and [forum](#). The RapidSmith project is merely a tool to make use of the XDL software technical interface and cannot be used without a valid and current license for the Xilinx ISE tools. The RapidSmith project is at the mercy of Xilinx in the availability of XDL and will attempt to accommodate updates and changes to the interface as they arise.

## Basic Syntax of XDL Files

XDL is a self-documenting file format in that each type of statement is generally preceded by comments that explain the syntax. Comments in XDL are denoted by using a '#' character at the beginning of a line. The '#' is also used in other constructs that are part of the language that do not fall on the beginning character of the line. In XDL there are four basic statements that make up the entire content of the file and description of the circuit.

### Design Statement

The design statement (represented as the `design.Design` class) is included in every XDL file (even hard macros) and there is only one design statement in a file. It includes global information such as the design name and part name of the targeted FPGA. It can also contain a list of attributes in a 'cfg' string. Below is an example of a design statement.

```
# =====
# The syntax for the design statement is:
# design <design_name> <part> <ncd version>;
# or
# design <design_name> <device> <package> <speed> <ncd_version>
# =====
design "designName" xc4vfx12ff668-10 v3.2 ,
    cfg "
        _DESIGN_PROP::BUS_INFO:4:OUTPUT:gpio<3:0>
        _DESIGN_PROP::PIN_INFO:gpio<0>:/top/PACKED/top/gpio<0>/PAD:OUTPUT:3:gpio<3\>:0>
        _DESIGN_PROP::PIN_INFO:gpio<1>:/top/PACKED/top/gpio<1>/PAD:OUTPUT:2:gpio<3\>:0>
        _DESIGN_PROP::PIN_INFO:gpio<2>:/top/PACKED/top/gpio<2>/PAD:OUTPUT:1:gpio<3\>:0>
        _DESIGN_PROP::PIN_INFO:gpio<3>:/top/PACKED/top/gpio<3>/PAD:OUTPUT:0:gpio<3\>:0>
        _DESIGN_PROP::PK_NGMTIMESTAMP:1231972339";
```

### Module Statement

Modules (represented as the `design.Module` class) are collections of instances and nets which can be described as hard macros if the instances are placed and nets are routed. A module will have a list of ports that determine the interface of the hard macro or module and each module will have its own list of instances and nets to describe the logic inside. An abbreviated module statement is shown below.

```
# =====
# The syntax for modules is:
#   module <name> <inst_name> ;
#   port <name> <inst_name> <inst_pin> ;
#   .
#   instance ... ;
#   .
#   net ... ;
#   .
#   endmodule <name> ;
# =====
module "moduleName" "anchorInstanceName" , cfg "_SYSTEM_MACRO::FALSE" ;
  port "portName1" "anchorInstanceName" "F2";
  port "portName2" "anotherInstanceInTheModule" "F4";
  ...
  inst "anchorInstanceName" "SLICEL", placed CLB_X14Y4 SLICE_X23Y8 , ...
  ...
  net "aNetInsideTheModule" , ...
  ...
endmodule "moduleName";
```

## Instance Statement

The instance statement (represented as the `design.Instance` class), which begins with the keyword ‘inst’, is an instance of an FPGA primitive which can be placed or unplaced depending if a tile and primitive site location are specified. The instance also has a primitive type (such as SLICEL, SLICEM, DCM\_ADV, ...). Instance names should be unique in a design to avoid problems in RapidSmith. Instances are configured with a ‘cfg’ string which is a list of attributes that define LUT content, and other functionality. An example of an instance statement is shown below.

```
inst "instanceName" "SLICEL",placed CLB_X14Y4 SLICE_X23Y8 ,
  cfg " BXINV::#OFF BYINV::#OFF CEINV::#OFF CLKINV::#OFF COUTUSED::#OFF CY0F::#OFF
  CY0G::#OFF CYINIT::#OFF DXMUX::#OFF DYMUX::#OFF F::#OFF F5USED::#OFF FFX::#OFF
  FFX_INIT_ATTR::#OFF FFX_SR_ATTR::#OFF FFY::#OFF FFY_INIT_ATTR::#OFF
  FFY_SR_ATTR::#OFF FXMUX::#OFF FXUSED::#OFF
  G:DCM_AUTOCALIBRATION_DCM_clock/DCM_clock/md/RSTOUT1:#LUT:D=A1
  _BEL_PROP::G:LIT_NON_USER_LOGIC:DCM_STANDBY GYMUX::#OFF REVUSED::#OFF SRINV::#OFF
  SYNC_ATTR::#OFF XBUSED::#OFF XMUXUSED::#OFF XUSED::#OFF YBUSED::#OFF
  YMUXUSED::#OFF YUSED::0 "
;
```

## Net Statement

The net statement (represented as the `design.Net` class) are the nets that describe inputs/outputs and routing of nets in a design. Nets can be of 3 different types: GND, VCC, or WIRE. The GND and VCC keyword must be present to mark a net as being sourced by ground or power, the keyword WIRE is not required as it is the default type. Nets also must have a unique name when compared with all other nets. Nets have two sub components to describe them: pins and PIPs. An example of a net statement is shown below.

Pins (represented as the `design.Pin` class) define the source and one or more sinks within the net. A pin is uniquely identified by the name of the instance where the pin resides as well as the internal name of the pin on this instance. It also has a direction of being an ‘outpin’ (source) or an ‘inpin’ (sink). A net can only have one source or ‘outpin’ in the net.

```

net "netName" ,
  outpin "instanceNameOfSourcePin" Y ,
  inpin "instanceNameOfSinkPin" RST ,
  pip CLB_X14Y4 Y_PINWIRE1 -> BEST_LOGIC_OUTS5_INT ,
  pip DCM_BOT_X15Y4 SR_B0_INT3 -> DCM_ADV_RST ,
  pip INT_X14Y4 BEST_LOGIC_OUTS5 -> OMUX8 ,
  pip INT_X15Y5 OMUX_EN8 -> N2BEG0 ,
  pip INT_X15Y7 N2END0 -> SR_B0 ,
;

```

PIPs (programmable interconnect points, represented by the `design.PIP` class) define routing resources used within the net to complete routing connections between the source and sinks. A PIP is uniquely described as existing in a tile (ex: `INT_X2Y3`) and two wires with a connection between them. Almost all PIPs are unidirectional (`'->'`) in that they can only go in one direction. Long lines are the one exception to that rule as they are bidirectional and are denoted by using a `'=='` symbol, however RapidSmith uses the `'->'` symbol for all PIPs as this does not cause the xdl converter any problems.

## Basic Syntax of XDLRC Files

XDLRC files are report files generated by the Xilinx `xdl` executable. During installation, RapidSmith will create XDLRC files and parse them for their pertinent information and then pack it into small device files that can be used later with the tool. Each construct found in XDLRC files and the corresponding RapidSmith representation is described in the remainder of this subsection.

### Tiles

```

# Example of an XDLRC tile declaration
  (tile 1 14 CLB_X6Y63 CLB 4
...
      (tile_summary CLB_X6Y63 CLB 122 403 148)
  )

```

Tiles (represented in the `device.Tile` class) are the building blocks of Xilinx FPGAs. Every FPGA is described as 2D array or grid of tiles laid out like a checker board (this can be seen also in the Part Tile Browser example). Each tile is declared with a `"(tile"` directive as shown above followed by the unique row and column index of where the tile fits into the grid of tiles found on the FPGA. The tile declaration also contains a name followed by a type with the final number being the number of primitive sites found within the tile. The tile ends with a `"tile_summary"` statement repeating the name and type with some other numbered statistics. Tiles can contain three different sub components, primitive sites, wires, and PIPs.

### Primitive Sites

```

# Example of an XDLRC primitive site declaration
  (primitive_site SLICE_X9Y127 SLICEL internal 27
    (pinwire BX input BX_PINWIRE3)
    (pinwire BY input BY_PINWIRE3)
    (pinwire CE input CE_PINWIRE3)
    ...
    (pinwire XMUX output XMUX_PINWIRE3)
  )

```

Primitive sites (represented in the `device.PrimitiveSite` class) are declared in tiles. A primitive site is a location on the FPGA that allows for an instance of that primitive type (primitive types are enumerated in the `device.PrimitiveType` enum) to reside. For example, in the declaration of a SLICEL primitive site above, any SLICEL instance can be placed at that site. A primitive site has a unique name (`SLICE_X9Y127`) and type (`SLICEL`). However, in some cases, more than one primitive type is compatible with a given primitive site. One example of this is the primitive type SLICEM (Virtex 4 slices that contain RAM functionality in the LUT among other enhancements to the SLICEL type) is a superset of SLICEL functionality. Therefore, a SLICEL primitive instance can be placed in a SLICEM primitive site. RapidSmith allows the developer to determine if a give site is compatible in the `device.PrimitiveSite` class using the method `isCompatiblePrimitiveType(PrimitiveType otherType)`.

Primitive site declarations in XDLRC also contain a list of pinwires which describe the name and direction of pins on the primitive site. The first pinwire declared in the example above is the BX input pin which is the internal name to the SLICEL primitive site. Pinwires have an external name as well to differentiate the multiple primitive sites that may be present in the same tile. In this case, BX of `SLICE_X9Y127` has the external name `BX_PINWIRE3`. RapidSmith provides mechanisms to translate between these two names in the `device.PrimitiveSite` class with the method `getExternalPinName(String internalName)`.

## Wire

```
# Example of an XDLRC wire declaration
(wire E2BEG0 5
    (conn CLB_X7Y63 CLB_E2BEG0)
    (conn INT_X8Y63 E2MID0)
    (conn CLB_X8Y63 CLB_E2MID0)
    (conn INT_X9Y63 E2END0)
    (conn INT_X9Y62 E2END_S0)
)
```

A wire as declared in XDLRC is a routing resource that exists in the tile that may have zero or more connections leaving the tile. In the example above, the wire `E2BEG0` connected to 5 other neighboring tiles. These connections (denoted by 'conn') are described using the unique tile name and wire name of that tile to denote connectivity. These connections are not programmable, but hard wired into the FPGA. Inter-tile connections are not programmable, however, intra-tile connections (PIPs, see below) are. RapidSmith must represent the routing resources of Xilinx FPGAs very carefully as a significant fraction of the FPGA description is routing. Therefore, the wire names (such as `E2BEG0`, ...) are enumerated into integers or Java primitive `int` data types using the `device.WireEnumerator` class. The `WireEnumerator` class keeps track of what integer value goes with each wire name and also for significant compaction of the FPGA routing description.

The wire connections are described using a relative tile offset to reuse data structure elements. The class used to represent these wires and corresponding connections is in the `device.Wire` class.

## PIP

```
# Example of an XDLRC PIP declaration
(pip INT_X7Y63 BEST_LOGIC_OUTS0 -> BYP_INT_B5)
```

A PIP (programmable interconnect point) is a possible connection that can be made between two wires. In the example above the PIP is declared in the tile and repeats the tile name for reference. It specifies two wires by name that both exist in that same tile (`BEST_LOGIC_OUTS0` and `BYP_INT_B5`) and declares that the wire `BEST_LOGIC_OUTS0` can drive the wire `BYP_INT_B5` if the PIP exists in a net's PIP list in a given design.

A collection of these PIPs in a net define how a net is routed and is consistent with saying that those PIPs are “turned on.” The connections are also represented in the `device.Wire` class as connections with a special flag denoting the connection as a PIP.

## Primitive Definitions

At the end of every XDLRC file (regardless of verbosity) contains a list of all primitive definitions for the Xilinx part. Primitive definitions are used mainly for reference and are reflected in the `primitiveDefs.*` package. In more recent Xilinx parts, some of the primitive definitions have been found to lack some information which may require special handling in RapidSmith. Currently, the primitive definitions are not widely used in RapidSmith.



# RAPIDSMITH STRUCTURE

This section details much of the complexity and theory behind the structure of RapidSmith. There are two main abstractions that developers need to be aware of; that of the device and design. A hierarchy of classes within RapidSmith can be seen in Figure 6 below.




Figure 6: (a) The classes involved in defining a design in RapidSmith, (b) The major classes involved representing a device.

## A RapidSmith Design

Designs in RapidSmith are represented and stored in the data structures found in the `design` package. The classes found in this package closely follow the constructs found in XDL design files to make the classes easier to follow and make the abstraction understandable to those who are familiar with XDL. For those who have less experience with XDL, see the [previous section on understanding XDL](#).

### Loading Designs

There is typically only one way to load a design with RapidSmith and that is to create a new design and call the method `loadXDLDesign(String fileName)` on that instance of the design. An NCD file can also be loaded indirectly by using methods in the `util.FileConverter` class which also conversion of an existing NCD file to XDL by calling the Xilinx `xdl` executable. Example code of how this could be done is shown below:

```
// Loading an XDL design
Design myDesign = new Design();
myDesign.loadXDLFile("myDesign.xdl");

// Loading an existing NCD file by converting it to XDL
String ncdFileName = "myOtherDesign.ncd";
String xdlFileName = FileConverter.convertNCD2XDL(ncdFileName);
if(xdlFileName == null){
    MessageGenerator.briefErrorAndExit("ERROR: Conversion of " +
        ncdFileName + " to XDL failed.");
}
Design otherDesign = new Design();
otherDesign.loadXDLFile(xdlFileName);
```

When loading a design, one must be conscious of everything that gets loaded. In RapidSmith, a [JavaCC](#)-based XDL parser (found in the `design.parser` package) reads and parses the given XDL file and populates the instance of the `Design` class respectively. When the parser populates the design with the design part name, it causes the corresponding device file of the Xilinx part as well as the wire enumerator to be loaded and populates the design. This is done by default because the primitive sites and wires loaded in with the design reference those same resources in the `device` class.

## Saving Designs

RapidSmith has a method also to save designs in the XDL format similar to the method for loading them. In a similar manner, the saved XDL file can be converted to NCD using the `FileConverter` class. Very little error checking is made when loading and saving XDL designs, but a good test would be the conversion to NCD as Xilinx runs several DRCs when the design is converted.

## A RapidSmith Device

A device is defined in RapidSmith as a unique Xilinx FPGA part that includes package information but not speed grade (such as xc4vfx12ff668). Each device contains specific information concerning its primitive sites, tiles, wires, IOBs, and PIPs that are available to realize designs. This information is made available through the Xilinx executable `xdl` in `-report` mode. See the [previous section](#) on XDLRC for more details on these device resources.

## Device

During the initial setup of RapidSmith, the Installer creates fully verbose XDLRC files (`$xdl -report -pips -all_conns <partName>`) for each device specified as command line parameters. After the creation of each XDLRC file, they are parsed, compacted by the Installer, and a device file is generated for later use. These device files are placed in

`$(RAPIDSMITH_PATH)/devices/familyName/partName_db.dat` and then the corresponding XDLRC file is deleted as they can be several gigabytes in size. These device files make accessing device information about a specific FPGA part much more convenient than a gigantic text file. Most of the device files are just a few megabytes or less and can be loaded in a few seconds or less. RapidSmith uses a custom form of serialization as well as a compression library to make sure the devices files are small and load quickly.

## Wire Enumerator

In order to make the device files small, each uniquely named wire is assigned to an integer as enumeration. This avoids moving strings around in memory which would be costly in terms of space and comparison times. RapidSmith has a class called `WireEnumerator` which enumerates all uniquely named wires in an FPGA family and has methods to convert to and from the wire name and enumeration or enum for short. It also stores information about each wire such as a direction or type which can be useful in building a router. Note that wires with the same name can occur several times within a device and they are uniquely identified not only by their name, but also by the tile in which they are present.

In order to create the wire enumeration files, a subset of XDLRC files must be parsed so that a complete set of wires can be enumerated. This is automatically done by the installer and the files are placed in `$(RAPIDSMITH_PATH)/devices/familyName/wireEnumerator.dat`. Only one wire enumerator is needed per FPGA family.

# Memory and Performance

Although Java typically has a reputation for being slow and a memory hog, RapidSmith has been able to create a very good device representation that is compact and fast loading, even for some of the largest parts offered by Xilinx. Performance and memory footprint figures are shown in the tables below for the device database file and data structure. These results were taken on Windows XP Professional SP3 (32-bit) with the Oracle (previously Sun) Java JVM build 1.6.0\_21-b07. The workstation used was an HP Compaq dc7800 CMT with an Intel Core 2 Duo 3.0 GHz (E6850) with 4GB RAM and 500GB SATA hard drive. Note that the `Device` and `WireEnumerator` classes are both required before loading any design in RapidSmith.

## Virtex 4 Device Performance and Memory Usage

Part Name	Heap Memory Footprint	File Size on Disk	Load Time from Disk
FX12	38.7 MB	599.0 KB	0.48 secs
FX20	63.4 MB	924.9 KB	0.79 secs
FX40	94.6 MB	1159.0 KB	1.02 secs
FX60	103.6 MB	1206.8 KB	1.42 secs
FX100	120.1 MB	1358.0 KB	1.51 secs
FX140	140.4 MB	1546.3 KB	1.63 secs
LX15	34.5 MB	231.8 KB	0.38 secs
LX25	37.3 MB	287.7 KB	0.46 secs
LX40	42.8 MB	348.5 KB	0.51 secs
LX60	64.0 MB	464.6 KB	0.80 secs
LX80	74.1 MB	596.7 KB	0.96 secs
LX100	74.2 MB	701.7 KB	0.97 secs
LX160	109.5 MB	875.6 KB	1.44 secs
LX200	115.5 MB	1010.6 KB	1.53 secs
SX25	62.8 MB	373.9 KB	0.76 secs
SX35	64.7 MB	441.0 KB	0.78 secs
SX55	74.8 MB	539.4 KB	0.93 secs

## Virtex 5 Device Performance and Memory Usage

Part Name	Heap Memory Footprint	File Size on Disk	Load Time from Disk
FX30T	62.9 MB	781 KB	0.86 secs
FX70T	83.9 MB	898.7 KB	1.03 secs
FX100T	107.6 MB	1014.3 KB	1.46 secs
FX130T	111.8 MB	1058.2 KB	1.5 secs
FX200T	129.3 MB	1227.3 KB	1.62 secs
LX20T	39.1 MB	497.4 KB	0.52 secs
LX30	41.4 MB	381 KB	0.53 secs
LX30T	46.1 MB	558.8 KB	0.56 secs
LX50	42.8 MB	417.4 KB	0.55 secs
LX50T	47.7 MB	586.7 KB	0.79 secs
LX85	71.6 MB	534.8 KB	0.98 secs

<b>LX85T</b>	77.4 MB	706.6 KB	1.01 secs
<b>LX110</b>	74.4 MB	588.9 KB	1.01 secs
<b>LX110T</b>	81 MB	761.9 KB	1.03 secs
<b>LX155</b>	102.6 MB	716.3 KB	1.43 secs
<b>LX155T</b>	113.5 MB	893.8 KB	1.47 secs
<b>LX220</b>	138.3 MB	862.4 KB	1.64 secs
<b>LX220T</b>	143.5 MB	1038.7 KB	1.87 secs
<b>LX330</b>	147 MB	1068.4 KB	1.88 secs
<b>LX330T</b>	152.9 MB	1250 KB	1.91 secs
<b>SX35T</b>	59.4 MB	596.5 KB	0.84 secs
<b>SX50T</b>	61.2 MB	630.1 KB	0.84 secs
<b>SX95T</b>	85.3 MB	754.4 KB	1.07 secs
<b>SX240T</b>	127.2 MB	1135.9 KB	1.64 secs
<b>TX150T</b>	89.9 MB	871.2 KB	1.1 secs
<b>TX240T</b>	122.4 MB	1111.2 KB	1.57 secs

### Virtex 6 Device Performance and Memory Usage

Part Name	Heap Memory Footprint	File Size on Disk	Load Time from Disk
<b>CX75T</b>	54 MB	583.9 KB	0.79 secs
<b>CX130T</b>	63.4 MB	714.8 KB	0.87 secs
<b>CX195T</b>	77.5 MB	853.2 KB	1.06 secs
<b>CX240T</b>	80.6 MB	937.3 KB	1.08 secs
<b>HX250T</b>	82.3 MB	973 KB	1.1 secs
<b>HX255T</b>	88.4 MB	1029.6 KB	1.12 secs
<b>HX380T</b>	97.8 MB	1287.3 KB	1.67 secs
<b>HX565T</b>	125.1 MB	1658.5 KB	1.81 secs
<b>LX75T</b>	54.1 MB	582.6 KB	0.79 secs
<b>LX130T</b>	63.3 MB	708.8 KB	0.87 secs
<b>LX195T</b>	77.3 MB	848.9 KB	1.07 secs
<b>LX240T</b>	80.6 MB	934.8 KB	1.08 secs
<b>LX365T</b>	107.2 MB	1186.4 KB	1.62 secs
<b>LX550T</b>	117.9 MB	1550.5 KB	1.8 secs
<b>LX760</b>	135.1 MB	1755.5 KB	2.48 secs
<b>SX315T</b>	106.5 MB	1157.1 KB	1.61 secs
<b>SX475T</b>	117.4 MB	1505.5 KB	1.79 secs

### Spartan 6 Device Performance and Memory Usage

Part Name	Heap Memory Footprint	File Size on Disk	Load Time from Disk
<b>LX100</b>	28.7 MB	530.1 KB	0.43 secs
<b>LX100T</b>	29.7 MB	620.1 KB	0.43 secs
<b>LX150</b>	34.6 MB	623.8 KB	0.52 secs

<b>LX150T</b>	35.7 MB	713.0 KB	0.52 secs
<b>LX16</b>	11.6 MB	234.8 KB	0.15 secs
<b>LX25</b>	20.2 MB	307.7 KB	0.28 secs
<b>LX25T</b>	21.1 MB	362.1 KB	0.26 secs
<b>LX45</b>	20.6 MB	359.2 KB	0.26 secs
<b>LX45T</b>	20.8 MB	421.8 KB	0.27 secs
<b>LX4</b>	13.4 MB	289.0 KB	0.20 secs
<b>LX75</b>	25.2 MB	468.1 KB	0.39 secs
<b>LX75T</b>	26.3 MB	555.6 KB	0.40 secs
<b>LX9</b>	9.0 MB	215.1 KB	0.12 secs

## Wire Enumerator Size and Performance

<b>Family Name</b>	<b>Heap Memory Footprint</b>	<b>File Size on Disk</b>	<b>Load Time from Disk</b>
<b>Virtex 4</b>	8.1 MB	234 KB	0.12 secs
<b>Virtex 5</b>	9.9 MB	264.4 KB	0.15 secs
<b>Virtex 6</b>	6.3 MB	171.3 KB	0.10 secs
<b>Spartan 6</b>	6.9 MB	168.0 KB	0.11 secs

# ROUTING IN RAPIDSMITH

This chapter is intended to help users and developers in understanding how routing resources are handled in RapidSmith. It also illustrates how to build on the existing classes to create custom routers.

## Wire Resources in RapidSmith

RapidSmith has a unique way of representing wires and connections for Xilinx devices. This approach was developed mainly to minimize disk and memory usage while also maintaining some level of efficiency and speed.

### Wire Representation

The wire enumerator class keeps a list of all uniquely XDLRC-named wires that exist in a given Xilinx FPGA family. Wires can span multiple tiles in the FPGA, however, the wire has a separate name for each tile in which it crosses. An example of this concept is illustrated in the DOUBLE lines found in several family architectures. A DOUBLE line is a wire that connects switch boxes either one or two hops away in a given direction. An example of this layout is given in Figure 7.




Figure 7 - A DOUBLE line in an FPGA illustrating how each part of the wire has a different name depending on which tile it is in.

In this example, we see a wire that can be driven by one point, E2BEG4, and can drive either E2MID4 in tile INT\_X2Y1 and/or E2END4 in tile INT\_X3Y1. However, the wire is assigned a name as it travels through the CLB tiles (CLB\_E2BEG4 and CLB\_E2MID4). For the purposes of RapidSmith, these wires have been removed from device files as they do not contribute to the overall possible connections a wire can make and simply add overhead to the device data structures. This technique has actually dramatically reduced the size of the devices files and improved routing speed as dead-end connections do not need to be examined.

In RapidSmith, these uniquely-named wire segments are represented either as a `String` or as an `int` or `Integer`. Often it is represented as an integer to save space and increase comparison speed with other wires. To illustrate how this representation works, here is some example Java code that exposes the wire segments:

```
// Load the appropriate Device and WireEnumerator
// (this is done automatically when loading XDL designs)
String partName = "xc4vfx12ff668";
Device dev = FileTools.loadDevice(partName);
WireEnumerator we = FileTools.loadWireEnumerator(partName);
// Here we pick a wire name
String wireName = "E2BEG4";
// Here we get the integer enum value for that wire name
int wire = we.getWireEnum(wireName);
// The wire enumerator also keeps information about these wire segments
```

```
// such as wire direction and type
WireDirection direction = we.getWireDirection(wire);
WireType type = we.getWireType(wire);
```

Now, there are actually several wires in an FPGA device with the same name. The wire E2BEG4 exists in almost every switch box tile in the FPGA. To uniquely identify routing resources in a device, a tile and its name or wire enumeration is required (that is, INT\_X1Y1 E2BEG4 is its unique representation).

In an effort to save space and ultimately reuse much of the routing connections, the `Wire` class is used to represent internal and external tile connections. Each tile has a hash map where the key is the integer enum value of the wire and the value is an array of `Wire` objects. Each `Wire` object contains the following information to define a connection:

```
/** The wire enumeration value of the wire to be connected to */
private int wire;
/** The tile row offset from the source wire's tile */
private int rowOffset;
/** The tile column offset from the source wire's tile */
private int columnOffset;
/** Does the source wire connected to this wire make a PIP? */
private boolean isPIP;
```

The `Wire` objects can define the connecting wire by using the integer enumeration value of the wire name and a relative offset of the tile differences between the two wires (again, relative to save space and increase reuse of the object). The `Wire` object also defines if the connection made is a programmable connection (or PIP).

When the row and column tile offsets are both 0, the connection exists within the same tile and is likely a PIP.

To query the connections that can be made from INT\_X1Y1 E2BEG4, here is some sample Java code to illustrate how this is done:

```
// Load the appropriate Device and WireEnumerator
// (this is done automatically when loading XDL designs)
String partName = "xc4vfx12ff668";
Device dev = FileTools.loadDevice(partName);
WireEnumerator we = FileTools.loadWireEnumerator(partName);
// Here we pick a wire name
String wireName = "E2BEG4";
// Here we get the integer enum value for that wire name
int wire = we.getWireEnum(wireName);
String tileName = "INT_X1Y1";
Tile tile = dev.getTile(tileName);
Wire[] wireConnections = tile.getWireConnections(wire);
for(Wire w : wireConnections){
    System.out.println(tileName + " " +
        wireName + " connects to " +
        dev.getTile(tile.getRow()-w.getRowOffset(),
            tile.getColumn()-w.getColumnOffset()) + " " +
        we.getWireName(w.getWire()) + " (is" +
        (w.isPIP()? " " : " not ") +
        "a PIP connection)");
}
```

#### Console Output:

```
INT_X1Y1 E2BEG4 connects to INT_X1Y1 BOUNCE1 (is a PIP connection)
INT_X1Y1 E2BEG4 connects to INT_X1Y1 BOUNCE2 (is a PIP connection)
```

```
INT_X1Y1 E2BEG4 connects to INT_X3Y1 E2END4 (is not a PIP connection)
INT_X1Y1 E2BEG4 connects to INT_X2Y1 E2MID4 (is not a PIP connection)
```

Routes in XDL are specified only with PIPs. Non-PIP connections (that is E2BEG4 to E2MID4, etc.) are not declared in an XDL Net since the connection is implied. The two wire segments are part of the same piece of metal on the FPGA. Thus, when declaring the routing resources used in a Net (the list of PIPs), these connections are not explicitly listed. However, the PIP connections are, for example:

```
net "main_00/i_ila/i_dt0/1/data_dly1_20" ,
  outpin "main_00/i_ila/i_dt0/1/data_dly1_20" XQ ,
  inpin "main_00/i_ila/i_yes_d/u_ila/idata_70" BY ,
  pip CLB_X16Y48 XQ_PINWIRE2 -> SECONDARY_LOGIC_OUTS2_INT ,
  pip CLB_X18Y48 BYP_INT_B4_INT -> BY_PINWIRE0 ,
  pip INT_X16Y48 SECONDARY_LOGIC_OUTS2 -> OMUX7 ,
  pip INT_X17Y48 OMUX_E7 -> E2BEG4 ,
  pip INT_X18Y48 E2MID4 -> BYP_INT_B4 ,
  ;
```

The listing of PIPs in XDL is arbitrary, that is, they do not always follow from one connection to the next.

## Basic Routing

RapidSmith has included an `AbstractRouter` class that allows for a common template so that routers can be constructed quite easily. However, the user should not feel restricted in using this template as it may not meet everyone's needs and/or requirements.

An example `BasicRouter` class has also been provided to illustrate how a router can be constructed easily. The `BasicRouter` class is ~400 lines of code. It is very simple and does not do any routing conflict resolution (it is a basic Maze router implementation) and it will commonly be unable to route certain connections in a design. Also, because the timing information for Xilinx parts is not publicly available, the router must use other means to optimize the router rather than delay. However, it does perform re-entrant routing, that is, it will attempt to route all nets that don't have any PIPs while keeping the original routed nets intact. If a net is impartially routed or improperly routed before given to the router, it does not resolve these problems. The behavior and mechanics of this router are described in the remainder of this section.

### Router Structure

The basic router provided in RapidSmith is based on a simple maze router algorithm. It does not allow routing resources to be used more than once, and thus, routing resources come on a first-come-first-served basis. This makes for a very simple implementation but does not resolve routing conflicts when they arise. The router chooses a route by iterating through a growing set of nodes, represented by the `Node` class. A node is a unique tile and wire combination to uniquely identify any routing wire available in the FPGA. Nodes are given a cost based on their Manhattan distance from the sink of the current connection to be routed and then placed in a priority queue. Those nodes with the smallest cost propagate to the bottom of the queue.

The least cost node of the queue is iteratively removed. With each removal, the node is examined for its expanding connections and those new potential nodes are also placed on the queue. Each time a node is removed, it is tested to see if it is the sink, if it is, the method traverses the path it has found and returns, otherwise it continues to expand more connections of the current node.

The router uses the following basic algorithm:

1. The central routing method, `routeDesign()` prepares the nets in the design for routing.



2. For each net in the design, `routeDesign()` will call `routeNet()`.
  - a. `routeNet()` prepares each inpin or sink in the net for routing.
    - i. If this is the first inpin of the net, it will only supply the outpin or source of the net as a starting point to the router.
    - ii. If this is the second or later inpin routed in the net, all intermediate points along those routes are added as starting points.
  - b. For each inpin, `routeNet()` will call `routeConnection()`.
    - i. `routeConnection()` initializes the priority queue of potential source nodes.
    - ii. `routeConnection()` calls the main routing method `route()` for each connection to be routed.
      1. The `route()` method iterates over the nodes in the priority queue, expanding their connections and adding new ones to the queue and putting more connections on the queue. The process continues until the sink is found.
3. After a net has been routed, the routing resources used will be marked as used to avoid reusing the resources twice.

## Routing Static Sources (VCC/GND)

One major preparation step in routing a full design is preparing where the static sources will be supplied from. The basic primitive in all Xilinx FPGAs to supply VCC and GND signals to a design is the TIEOFF. The TIEOFF accompanies every switch matrix and has several connections to all sink connections to its neighboring logic tile (CLB, BRAM, DSP, etc.). It has 3 pins, HARD0 or GND, KEEP1 (VCC) and HARD1 (VCC). By default, without any configuration, it seems that pins will default to KEEP1. Some pins, however, require a HARD1 when specified to be driven with VCC.

The `StaticSourceHandler` class takes care of partitioning the various nets and sinks into their respective tiles and instantiating the TIEOFF automatically. It also will instance SLICES when necessary. It also “reserves” certain routing resources for certain nets that could potentially introduce routing conflicts later. These reserved nodes are released just before the net is routed in the basic router.

## Routing Clocks

When routing clocks, it is quite important that they get routed to the appropriate clock tree routing resources. The best current method to determine this is based on the `WireDirection` (the type CLK was placed in `WireDirection` because there are certain CLK wires that also fell into certain `WireType` categories). The cost function for determining node position in the priority queue take into account clock wires and significantly reduces their cost when routing clock nets.

# APPENDIX

Here is just a grouping of useful topics that may not fit in the rest of this document.

## Appendix A: Modifying LUT Content

LUTs (look up tables) found in Xilinx slices can be easily modified using RapidSmith. LUT content is stored in an attribute in an instance of a SLICEL or SLICEM or whatever type of SLICE the device has. Often the name of the LUT is a single letter such as F or G as in the Virtex 4 family. Virtex 5 FPGAs have 4 LUTs in a slice and are called “{A, B, C, D} {5, 6} LUT” which have the capability to act as a 5 input or 6 input LUT.

### LUT Equation Syntax

Xilinx uses the following syntax for operators in a LUT equations string:

Operator	Operator Meaning
*	Logical AND
+	Logical OR
@	Logical XOR
~	Unary NOT

The parenthesis characters are also used to denote precedence in equation. Valid equation values are A1, A2, A3 and A4 for 4-input LUTs with an additional A5 for 5-input LUTs and A6 for 6-input LUTs. Some examples of LUT equations are:

LUT Equation	Equation Meaning
A1	Passes through the signal A1
A1*A2	A 2-input AND gate using A1 and A2
~A4	Inverts the signal on A4
(A4+ (A1+ (A2+A3) ) )	A multi-level OR of 4 inputs

### XDL LUT Equation Syntax

An instance is configured by zero or more attributes in a list where an attribute has the pattern

```
<Physical Name>:<Optional Logical Name>:<Value>.
```

The physical name of a LUT has been mentioned at the beginning of this subsection as being F or G for Virtex 4 parts and “{A, B, C, D} {5, 6} LUT” for Virtex 5 parts. The optional logical name is generally chosen by the synthesis tools or whatever name propagated through to Xilinx NGDBuild and Map when the signal was converted. It is only for reference back to the user’s original design and should correspond with the output of the LUT. The value has the following syntax:

```
#LUT:D=<equation>
```

This syntax can also be found in the XDLRC primitive\_defs. If we put this altogether with our examples above we would get:

```
F:mySignal0:#LUT:D=A1
G:mySignal1:#LUT:D=A1*A2
F:mySignal2:#LUT:D=~A4
G:mySignal3:#LUT:D=(A4+(A1+(A2+A3) ) )
```

Notice that the value also contains the colon (':'). RapidSmith defines the separation of the three components of an attribute as the first and second colons, any colons found after the second colon is part of the attribute value.

It is also of some value to recognize that the router can re-arrange LUT inputs to make routing easier and this can change equation to some extent.

## ***Appendix B: Hard Macros in XDL and RapidSmith***

A hard macro is a collection of primitive instances and nets that have been placed and possibly routed. The XDL construct that represents a hard macro is the “module”. A hard macro file only has a single module and no other nets, instances or module instances present in the file. The design name is “\_\_XILINX\_NMC\_MACRO” which is a static variable in the Design class:

```
public static final String hardMacroDesignName = "__XILINX_NMC_MACRO";
```

### **Xilinx NMC files**

The NCD file format is not used for hard macros; instead, Xilinx uses the NMC format which is fairly similar. NMC files can be converted to and from XDL using the same xdl executable. For example:

```
xdl -ncd2ndl myHardMacro.nmc [optional: myHardMacro.xdl]
xdl -xdl2ncd myHardMacro.xdl myHardMacro.nmc
```

The user must specify the output file name for NMC files when converting from XDL to NMC, otherwise it will output a file with the same name but NCD extension and the Xilinx tools won't read it properly. You may also need to apply the `-force` option if there is an error.

### **RapidSmith Hard Macro Generator**

RapidSmith provides a generic hard macro generator that will take a placed and routed NCD design file (or XDL equivalent) and replace IOBs with module ports among other transformations to make it a valid and legal hard macro. The hard macro generator provided is the same as used in the published paper, “Using Hard Macros to Reduce FPGA Compilation Time”, C. Lavin, et al., FPL'2010. This class is found in the `util` package and can be run from the command line:

```
java edu.byu.ece.rapidSmith.util.HardMacroGenerator <input.xdl |
input.ncd> <output file type: xdl | nmc> [optional: original_toplevel.vhdl]
```

The hard macro generator has to remove TIEOFFs (which supply GND and VCC) because they are not legal primitives in hard macros. Therefore, the hard macro generator creates input ports accordingly. It also renames ports according to primitive instance names of the IOBs and will attempt to name them as they were found in the original HDL. However, it only can guess, and for a more complete naming, the top level VHDL file can optionally be provided as a parameter to help the generator get the port names correctly. Unfortunately, the generator cannot parse Verilog.